GoalOS-native alpha-AGI Ascension using AGIALPHA

Sovereign RSI Architecture for Evidence-Governed Intelligence Organizations

A model can answer. An agent can act. A sovereign intelligence institution must prove.

Vincent Boucher

President, QUEBEC.AI & MONTREAL.AI | Sovereign AI, AI-First Institutions, AI Governance & Governed Agents

Corporate Sovereign RSI Edition v6.1

Institutional boundary

This paper defines a sovereign RSI architecture, proof standard and implementation direction. It does not authorize deployment or claim achieved AGI, superintelligence, guaranteed economic return, legal approval, tax approval, security approval or Kardashev Type II achievement.

Reader Map

This edition is designed for board members, sovereign AI leaders, protocol engineers, commercialization teams and reviewers. It uses a ledger-inspired ivory, navy and gold visual language with no floating text boxes. Every figure is bounded; every table is page-safe; every claim is evidence-gated.

Executive thesis

GoalOS-native alpha-AGI Ascension turns the original agent-market vision into a proof-governed sovereign intelligence operating system. GoalOS defines what should improve. AGIALPHA coordinates proof-settled work. Evidence Dockets make claims auditable. The Proof Gradient decides what may influence future work.

Abstract

AI-first institutions should not merely automate work; they should generate evidence sufficient to govern what may improve next. This paper presents GoalOS-native alpha-AGI Ascension using AGIALPHA: a sovereign RSI architecture in which institutional aims become commitments, agent work becomes ProofBundles, reviewers produce attestations, the Proof Gradient selects only evidence-backed upgrades, and Evidence Dockets make claims auditable. AGIALPHA is treated as a utility coordination asset for proof-settled work: mission posting, bonds, reviewer accountability, settlement, proof-card actions, credentials, reputation and access. The paper does not claim achieved AGI, superintelligence, guaranteed financial return or deployment authorization. It defines the institutional machinery required to test whether useful AI workflow work can become governed, reusable, compounding capability.

Canonical Contributions

Contribution	Why it matters
Sovereign RSI constitution	RSI is defined as proof-backed upgrade rights, not uncontrolled self-modification.
GoalOS-native reimplementation	The original Ascension primitives become GoalOS, AEP-001, Evidence Dockets, Proof Cards, credentials and reputation.
AGIALPHA proof-settlement model	AGIALPHA coordinates proof work without exposing private intelligence or creating investment claims.
Large multi-agent operating theater	Specialist roles cooperate through bounded authority and artifact discipline.
Evidence Docket standard	Consequential claims are backed by commitments, runs, proofs, evals, risk, settlement, rollout and rollback evidence.
Sovereign commercial deployment path	Self-serve products, proof rooms, sponsor missions and enterprise pilots become one compounding system.
Falsifiable proof program	The architecture becomes credible only through real tasks, baselines, replay, attestations, delayed outcomes and independent review.

1. Sovereign RSI Architecture

The architecture is not a swarm and not a design metaphor. It is an operating theater in which every role has bounded authority and every transition emits evidence. RSI means that accepted artifacts earn limited rights to shape future missions.

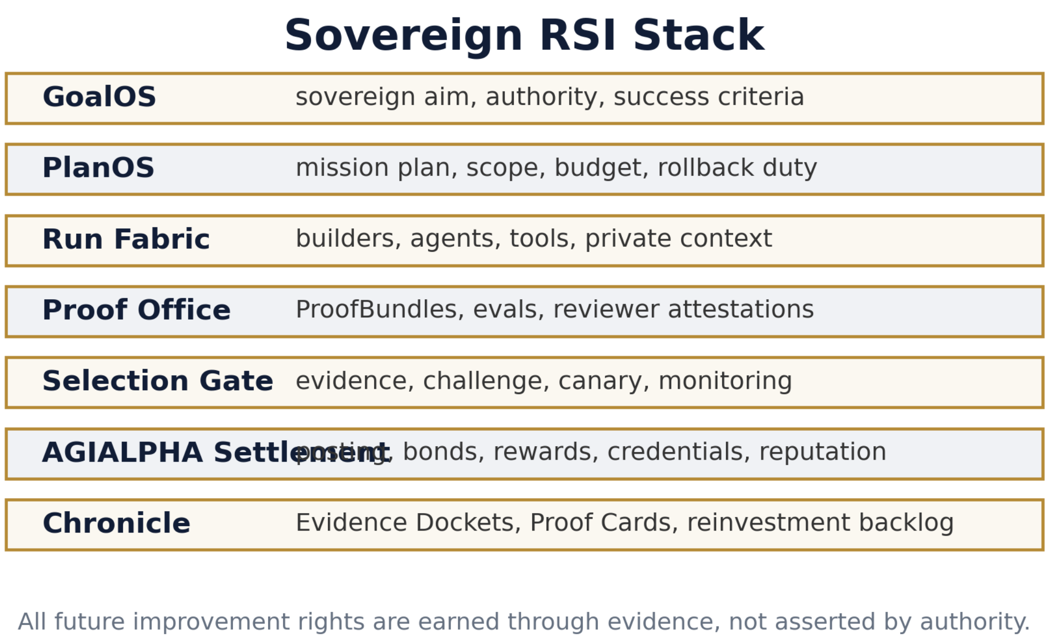

Figure 1. GoalOS-native sovereign RSI stack.

2. Reimplementation Mapping

The reimplementation preserves the durable primitives of the original Ascension concept while replacing speculative or ambiguous mechanics with auditable institutional objects.

Original primitive	GoalOS-native primitive	Upgrade
Ascension thesis	GoalOS-native proof-of-evolution thesis	From visionary agent constellation to evidence-governed intelligence organization.
Insight	GoalOS Opportunity Intelligence	Workflow opportunities replace unsupported foresight.
Nova-Seeds	Proof Seeds / Opportunity Dockets	Mission-ready artifacts with evidence requirements and claim boundaries.
FusionPlans	PlanOS Mission Plans	Executable routes with budget, scope, risk and rollback duties.
MARK	AGIALPHA Proof Mission Board	A market for proof work rather than speculative foresight.
Sovereign	Proof Room / Enterprise RSI Pilot	Institutional transformation through auditable workflow improvement.
Marketplace	AGIALPHA Proof Work Network	Sponsor-builder-reviewer coordination

Original primitive	GoalOS-native primitive	Upgrade
		with credentials and reputation.
Agents	Builders / Agent Operators	Accountable operators who generate evidence.
Jobs	GoalOS Proof Missions	Goal, metric, bounty, proof requirements, reviewer rubric and public-safe output.
Validator Council	Reviewer Network + Selection Gate	Bonded review, challenge windows, scope control and rollback readiness.
Architect	Proof Gradient / AEP-001	A governed evolution law for what may improve future work.

3. AGIALPHA Proof-Settlement Chain

AGIALPHA is used as a utility coordination asset for proof-settled work: mission posting, builder bonds, proof bonds, reviewer bonds, settlement, proof-card actions, credentials and reputation. Normal product sales do not need token friction.

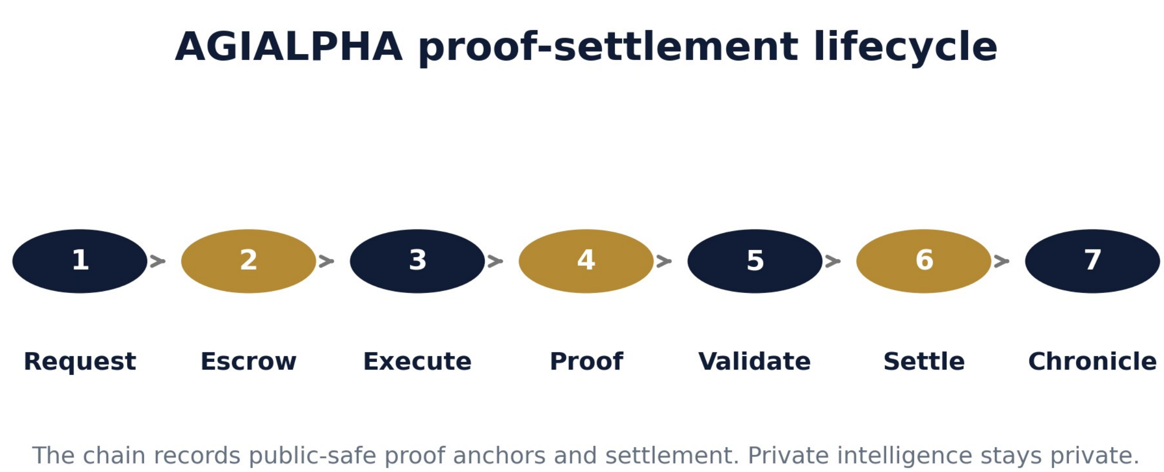

Figure 2. Request -> Escrow -> Execute -> Proof -> Validate -> Settle -> Chronicle.

4. Multi-Agent Coordination Fabric

The system demonstrates large-scale coordination through specialization. Each role emits a concrete artifact; each artifact can be reviewed, challenged, accepted, rejected or rolled back.

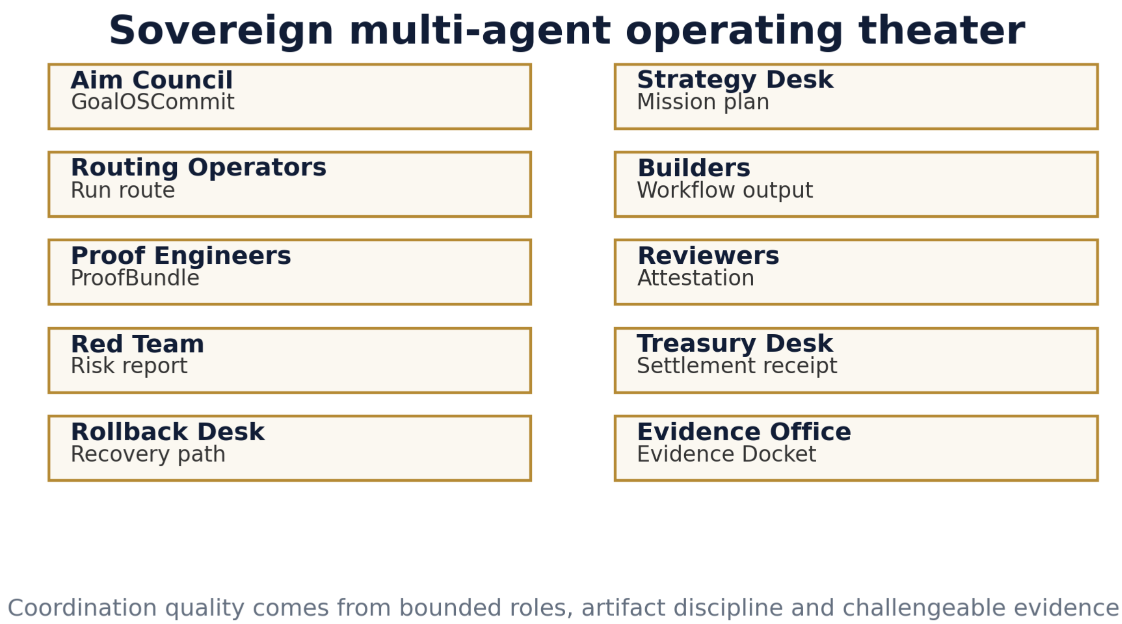

Figure 3. Bounded specialist-agent operating theater.

5. RSI as Operational Law

The system does not rely on open-ended self-modification as a promise. RSI is implemented as a law of propagation: no artifact can influence future work until it survives evidence, evaluation, reviewer validation, challenge, canary scope, monitoring and rollback readiness.

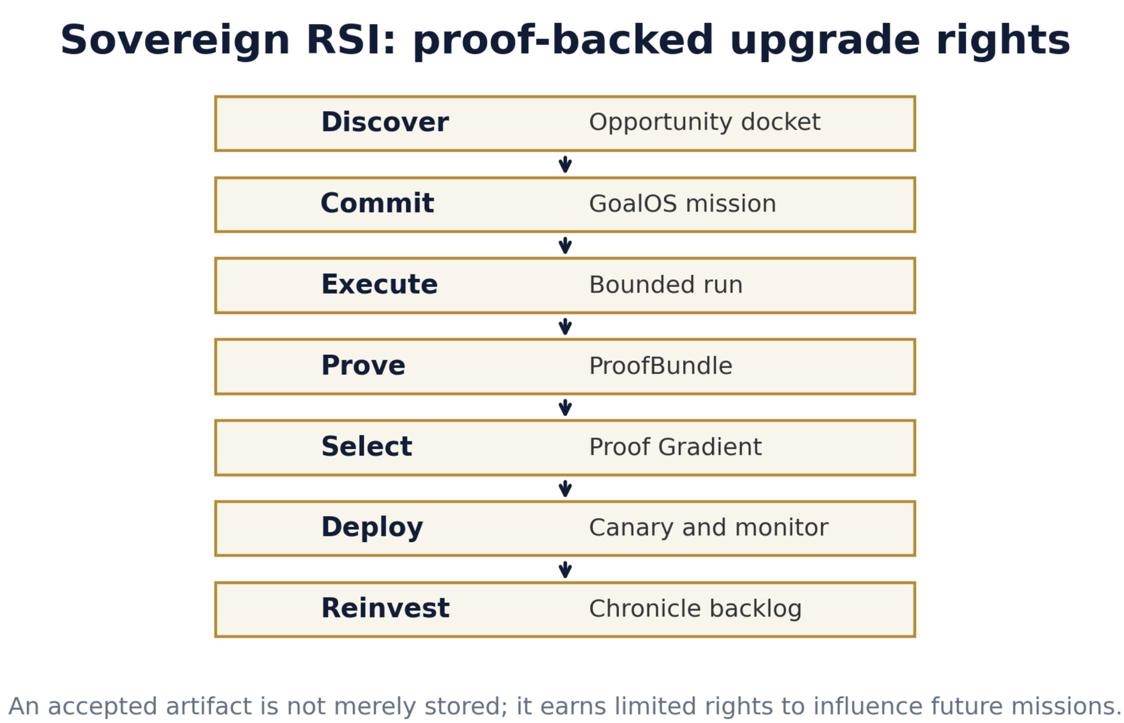

Figure 4. Proof-backed upgrade rights.

Sovereign RSI law

An artifact may shape future work only after evidence, evaluation, reviewer validation, scope control, challenge window, canary rollout, monitoring and rollback readiness.

6. Goals Used

Goal	Explanation
Sovereign Aim Integrity	Every mission preserves declared aim, risk class and authority boundary.
Proof-Settled Work	Work becomes reusable only after evidence, evaluation and review.
RSI With Rollback	Accepted improvements include monitoring and recovery paths.
Private Intelligence, Public Proof	Sensitive work stays off-chain; public-safe anchors become verifiable.
Reputation Through Evidence	Reputation follows approved work and review quality.
Commercial Compounding	Proof Cards and Evidence Dockets create trust that can sell the next GoalOS tier.
Capacity Reinvestment	Verified capacity gains fund better agents, evals, safety and useful infrastructure.

7. Plans Used

Plan	How it is used
Opportunity Docket Plan	Convert workflow pain into a sponsor-ready mission with a measurable success condition.
Mission Settlement Plan	Use AGIALPHA for posting, claim bonds, proof bonds, reviewer bonds and reward settlement.
Evidence Docket Plan	Package commitments, ProofBundles, eval attestations, risk records and settlement receipts.
Selection Gate Plan	Promote only artifacts that pass evidence, evals, risk, scope, challenge and rollback gates.
Sovereign RSI Plan	Turn accepted improvements into versioned artifacts that can shape future missions.
Season 001 Plan	Launch five focused proof missions before scaling the mission board.
Deployment Gate Plan	Progress from repository evidence to rehearsal, external review and formal authorization.

8. Skills Used

Skill	Operational role
GoalOS Commit Compiler	Turns sovereign intent into a signed, versioned mission object.
Opportunity Docket Writer	Frames a workflow opportunity with pain, metric, evidence and sponsor value.
Mission Router	Routes proof missions to qualified builders and reviewers.
Builder Operator	Executes workflow improvement and emits traceable artifacts.
ProofBundle Assembler	Packages outputs, hashes, scorecards, cost, latency and replay material.
Reviewer Attestation	Validates whether evidence supports the claimed improvement.
Red Team Examiner	Searches for unsafe claims, private data leakage and brittle improvements.
Proof Gradient Judge	Applies selection criteria to decide what may influence future work.
Rollout Planner	Defines scope, canary, monitoring and rollback path.
Evidence Docket Publisher	Creates the proof room used by executives, auditors and partners.
Settlement Steward	Coordinates AGIALPHA reward, bonds, credentials and reputation updates.
Chronicle Curator	Stores approved proof memory for future mission design.

9. Evidence Docket as Proof Room

The Evidence Docket is the institutional proof room. It is designed to let executives, auditors, reviewers and sponsors inspect what was claimed, what was done, what evidence exists, who reviewed it, how it was settled and whether rollback remains possible.

Evidence Docket contents

Commit evidence, run evidence, ProofBundle, eval attestation, risk ledger, settlement receipt, rollout scope, rollback path and public claim boundary.

10. AGIALPHA Settlement, Credentials and Reputation

Mechanism	Institutional effect
Mission posting	Sponsors commit resources to defined proof work.
Builder bond	Builders accept accountability for completion.
Proof bond	Proof submitters stake on the seriousness of the evidence.
Reviewer bond	Reviewers have skin in the quality of their validation.

Mechanism	Institutional effect
Proof Card	Public-safe proof artifact creates trust without leaking private intelligence.
Credential	Accepted work becomes legible capability.
Reputation	Future access is guided by evidence, not assertion.

11. Selection Law

A candidate upgrade is selected only if verified improvement exceeds its cost, risk and novelty penalties, and only if all hard gates pass.

Selection condition

Promote(a) is TRUE only when Proof(a) is sufficient, Eval(a) passes, Risk(a) is acceptable, Scope(a) is authorized, Challenge(a) clears, Canary(a) passes, Monitor(a) is active, and Rollback(a) is ready. Otherwise: revise, reject, quarantine or retire.

12. Capacity Reinvestment Horizon

The civilization-scale horizon is treated as a disciplined research program. The measurable near-term pathway is: work creates proof; proof creates trusted artifacts; trusted artifacts improve future work; improvements compound capacity; capacity can be reinvested into compute, energy efficiency, safety, public-good infrastructure and scientific or institutional capability.

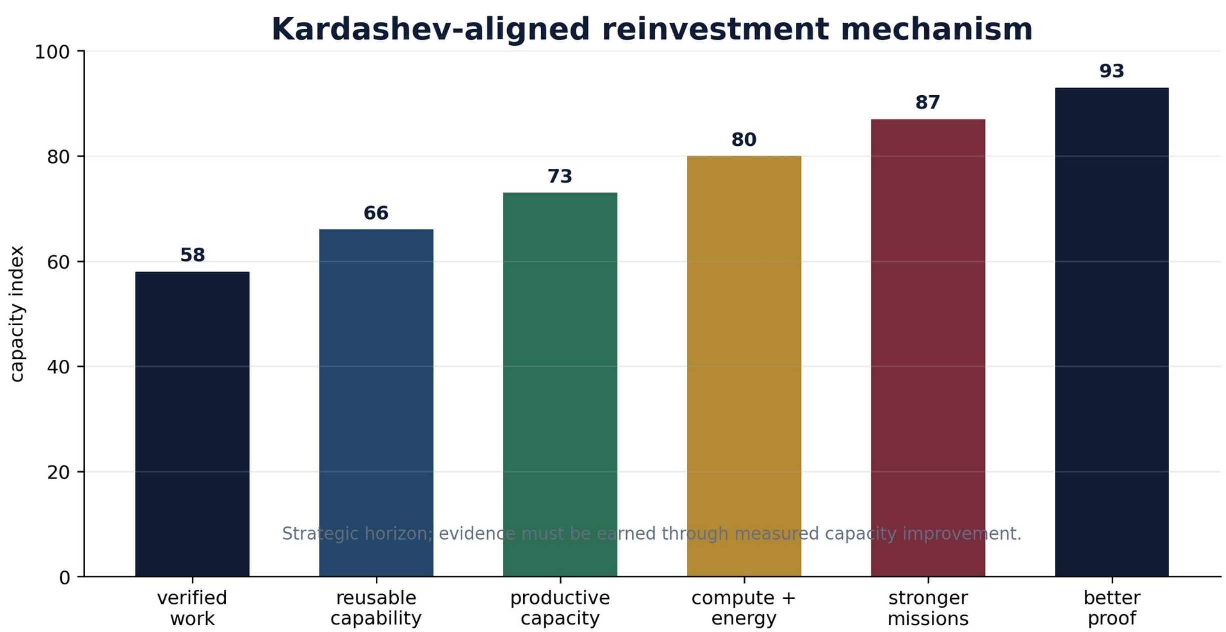

Figure 5. Clean reinvestment mechanism; no present-tense endpoint claim.

13. Readiness Ladder

The repository and implementation package can be strong before deployment authority exists. Authority must be earned through compile/tests, rehearsal, Evidence Docket, external review, legal/tax/security review, governance ceremony, public claims review and formal approval.

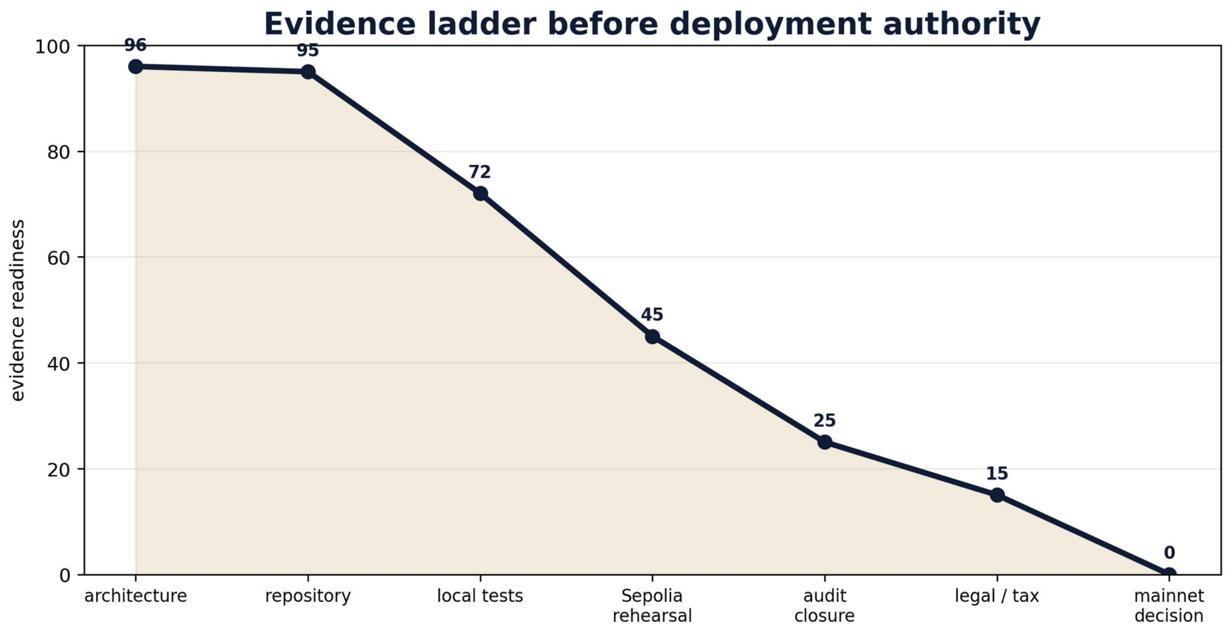

Figure 6. Evidence gates before deployment authority.

14. Season 001 Proof Mission Program

Mission	Workflow	Proof artifact	Business value
PM-001	Customer Support Reply Workflow	Proof Card 001	Support quality and safe escalation
PM-002	Claims-Safe Launch Content	Proof Card 002	Public messaging and claim boundary
PM-003	Workshop Application Qualification	Proof Card 003	Pipeline quality and conversion
PM-004	Sponsor Proof Job Creation	Proof Card 004	Sponsor demand formation
PM-005	Builder / Reviewer Quality	Proof Card 005	Proof economy reliability

15. Sovereign Domain Applicability

The architecture is most useful where private intelligence must stay protected while evidence, review, reputation and rollout decisions must become institutionally verifiable.

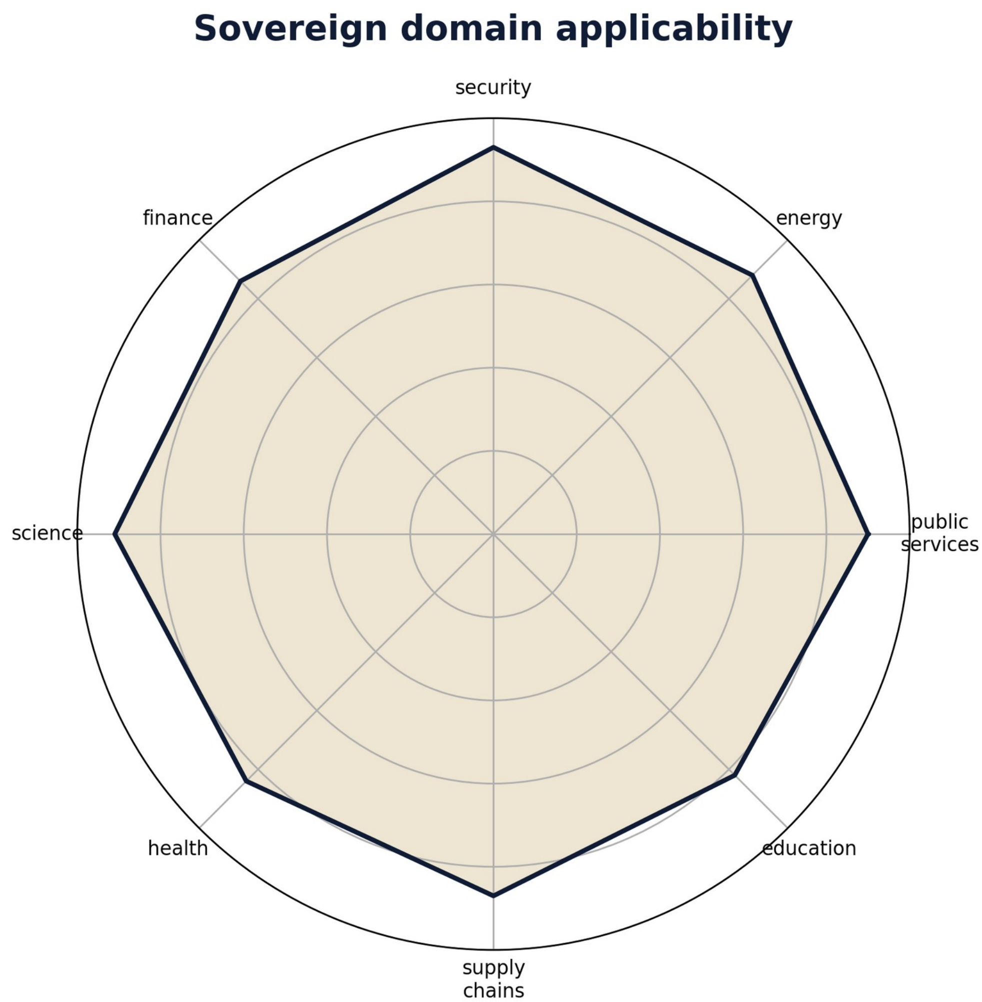

Figure 7. Sovereign domains where proof-governed RSI is most useful.

16. Contract and Registry Architecture

Module	Purpose
Mission Board	Posts missions, handles claims, proof submission, reviewer validation, settlement and states.
Proof Card Registry	Registers public-safe proof-card hashes and URIs.
Proof Credential	Issues non-transferable credentials for accepted proof work.
Reputation Ledger	Tracks evidence-backed non-transferable reputation.
Launch Gate Registry	Records authorization evidence and blocks premature deployment claims.
Legacy Evidence Registry	Preserves AGIJobManager as genesis context without making it the current architecture.
Evidence Docket Schema	Defines what a proof room must contain before claims mature.

Current status: evidence-ready audit candidate. Not external audit closure. Not legal approval. Not tax approval. Not security approval. Not deployment authorization.

17. Falsification Tests

Failure mode	What would falsify the claim
Proof fails to predict improvement	Accepted artifacts do not outperform baselines.
Reviewers are gameable	Low-quality or collusive reviews pass selection.
Cost exceeds value	Coordination overhead exceeds verified value created.
Rollback fails	Accepted improvements cannot be reversed safely.
Privacy boundary fails	Private evidence leaks through public proof anchors.
AGIALPHA use is narrative-only	Token activity does not coordinate real posting, bonding, review or settlement.
Evidence Dockets are incomplete	Claims cannot be replayed or inspected independently.

18. Claim Boundary

This paper defines a research architecture, implementation direction and evidence program. It does not authorize deployment. It does not guarantee economic outcomes. It does not claim achieved AGI, ASI, superintelligence, autonomous sovereignty, investment return, audit closure, legal approval, tax approval, security approval or Kardashev Type II achievement.

19. Canonical Constitution

GoalOS-native alpha-AGI Ascension

GoalOS sets the aim. AGIALPHA coordinates proof-settled work. AEP-001 defines valid evidence. The Proof Gradient decides what may evolve. The Artifact Vault stores reusable intelligence. The Run Fabric executes bounded agents. The Proof Ledger records what happened. The Selection Gate promotes only what survived evidence. The Chronicle preserves institutional memory. The intelligence stays private. The proof becomes verifiable. The network improves only through verified work.

References

Boucher, Vincent. GoalOS: The Proof-of-Evolution Constitution. AEP-001 Institutional Standard. QUEBEC.AI and MONTREAL.AI, June 2026.

Boucher, Vincent. AGI ALPHA: A Scalable Substrate for Intelligence Organizations. April 2026.

AGIALPHA token tracker, Ethereum mainnet: 0xA61a3B3a130a9c20768EEBF97E21515A6046a1fA.

MontrealAI/goalos-agialpha-ascension repository foundation and evidence-ready implementation packages.