

GoalOS-Native alpha-AGI Ascension using AGIALPHA

A Sovereign RSI Architecture for Evidence-Governed Intelligence
Organizations

By Vincent Boucher

President, QUEBEC.AI & MONTREAL.AI

Sovereign AI | AI-First Institutions | AI Governance | Governed Agents

Core thesis

**GoalOS sets sovereign aims. AGIALPHA coordinates proof-settled work.
AEP-001 defines valid evidence. The Proof Gradient decides what may evolve.
The intelligence stays private. The proof becomes verifiable.**

This paper defines a practical sovereign RSI architecture: improvements earn the right to propagate only through evidence, evaluation, reviewer validation, scope control, monitoring, and rollback readiness.

Sovereign RSI Edition v6.3

Abstract


AI systems should not merely act, remember, or self-improve. They should produce evidence strong enough to govern what may improve next. GoalOS-native alpha-AGI Ascension using AGIALPHA defines a sovereign RSI architecture in which institutional aims become signed commitments, bounded agents execute proof missions, AGIALPHA coordinates proof-settled work, reviewers validate evidence, Evidence Dockets make claims auditable, and the Selection Gate grants propagation rights only to verified improvements.

The paper reinterprets alpha-AGI Ascension as a practical proof-work economy: opportunities become proof missions; builders produce replayable ProofBundles; reviewers issue attestations; Proof Cards create public-safe trust; credentials and reputation route future work; and accepted artifacts become proof-backed upgrade rights. The result is not an ungated swarm or a speculative foresight market. It is an evidence-governed intelligence institution whose recursive self-improvement occurs through proofs, baselines, challenge windows, canaries, rollbacks, and durable chronicle memory.

Claim boundary: this is a protocol and architecture paper. It does not claim achieved AGI, ASI, superintelligence, guaranteed economic return, legal approval, tax approval, security approval, energy abundance, or Kardashev Type II achievement. It defines the machinery required to test whether verified work can become reusable capability and whether reusable capability can compound under sovereign governance.

1. Executive Thesis

The strategic advantage of GoalOS-native alpha-AGI Ascension is not raw autonomy. It is governed improvement. Frontier models make action cheap; they do not make institutional trust cheap. This architecture makes trust programmable by requiring every material improvement to pass through evidence, evaluation, reviewer validation, scope control, canary rollout, monitoring, and rollback readiness before it can affect future work.



2. Canonical Contributions

Contribution	Definition	Why it matters
Sovereign RSI	Defines recursive self-improvement as proof-backed upgrade rights rather than ungated self-modification.	Practical governance of improvement.
GoalOS Aim Layer	Converts strategic intent into signed commitments with authority, risk, budget, evaluators, and rollback.	Prevents local agent optimization from drifting away from institutional goals.
AGIALPHA Settlement Rail	Coordinates posting, bonds, proof submission, reviewer validation, settlement, and reputation.	Turns work into replayable proof-settled machine labor.
Evidence Docket	Packages claims, baselines, proof bundles, costs, risks, replay paths, and challenge records.	Makes public claims auditable rather than promotional.
Selection Gate	Promotes only candidates with proof, evaluations, scope controls, canary evidence, monitoring, and rollback readiness.	Separates capability from authority to propagate.
Capacity Reinvestment	Channels verified gains into compute, data, tooling, validators, security, and future mission capacity.	Connects near-term workflow proof to the long strategic horizon.

3. Reimplementation Mapping

The reimplementation preserves the ambition of alpha-AGI Ascension while replacing ambiguous autonomy with auditable proof mechanics. It turns foresight, seeds, markets, agents, validators, and sovereign transformation into concrete operating primitives: Opportunity Dockets, Proof Missions, Proof Cards, Credentials, Reputation, Evidence Dockets, and Selection Certificates.

Original layer	GoalOS-native layer	Upgrade
alpha-AGI Insight	GoalOS Opportunity Intelligence	Turns foresight into auditable workflow opportunity dockets.
Nova-Seeds	Proof Seeds	Replaces speculative seed language with mission-ready, public-safe opportunity artifacts.
MARK	AGIALPHA Proof Mission Board	Moves from futures on foresight to sponsor-funded proof missions.
Sovereign	GoalOS Proof Room / Sovereign Pilot	Transforms enterprises through governed workflow improvement, evidence, rollout, and rollback.
Marketplace	AGIALPHA Proof Work Network	Coordinates sponsors, builders, reviewers, credentials, reputation, and settlement.
Agents	Builders / Operators	Executes bounded workflows and emits proof bundles.
Validator Council	Reviewer Network + Selection Gate	Validates evidence through rubrics, challenge windows, and reputation.
Architect	Proof Gradient + Agent Evolution Protocol	Turns meta-optimization into governed RSI.

4. AGIALPHA Proof-Settlement Lifecycle

AGIALPHA is used as the existing coordination asset for proof-settled work. Its institutional role is utility coordination: mission posting, bonds, settlement, metrology, reviewer coordination, and protocol operations. It is not described as equity, revenue share, yield, ownership, or guaranteed return.

AGIALPHA Proof-Settlement Chain




Every settlement requires replayable proof. Every accepted proof becomes chronicle memory.
No ProofBundle, no settlement. No replay, no authority. No rollback, no propagation.

5. Sovereign RSI: The Practical Mechanism


The RSI mechanism is institutional, not magical. A model does not freely rewrite itself. Instead, the intelligence organization recursively improves workflows, proof templates, validator rubrics, tools, mission environments, agent coalitions, archives, and allocation policies. The object of improvement is the proof-carrying operating system around the agents.

The RSI loop is: detect a gradient; generate candidates; execute bounded work; package proof; validate; settle; archive; allocate; monitor delayed outcomes; and promote only when rollback is ready. High novelty creates higher scrutiny, not automatic authority.



6. Large Multi-Agent Coordination Fabric

The system is designed as an operating theater of specialized agents and offices, not a single autonomous actor. Each role has bounded authority, emits a named artifact, and participates in evidence creation. The goal is not to maximize action volume; the goal is to maximize verified useful work under cost, risk, and rollback constraints.



7. Skills Used

Skills are not vague capabilities. They are named operational competencies with inputs, outputs, and verification obligations.

Skill	Layer	What it does
GoalOSCommit Compiler	Aim	Turns sovereign intent into signed mission commitments with risk, budget, evaluators, and rollback.
Mission Foundry	Planning	Generates proof missions, environments, validators, and curricula for high-value workflow domains.
Constellation Router	Coordination	Routes work to builder coalitions, proof engineers, reviewers, red teams, and rollback desks.
ProofBundle Packager	Evidence	Packages traces, hashes, telemetry, signatures, costs, risks, replay data, and validation inputs.
Reviewer Council Orchestrator	Validation	Coordinates bonded reviewers, commit-reveal logic, dispute windows, and slashing paths.
Selection Gate Judge	Evolution	Awards propagation rights only after proof, evals, canary, monitoring, and rollback readiness.
Capacity Allocator	Reinvestment	Converts verified gains into stronger future compute, data, tools, security, and mission capacity.

8. Plans Used

Plan	Operational motion	Proof output
Season 001 Proof Missions	Begin with five high-signal missions: support workflow, claims-safe launch, workshop qualification, sponsor missions, reviewer quality.	Create first public-safe proof cards and reviewer history.
Sovereign Proof Room Pilot	Select one department workflow, baseline it, run builders, validate, publish a docket, and expand only after rollback readiness.	Makes enterprise adoption auditable.
AGIALPHA Settlement Rehearsal	Use test environment to run post -> claim -> proof -> validate -> proof card -> credential -> reputation -> settlement.	Proves the proof-work economy before mainnet authorization.
Evidence Docket Publication	Publish public-safe claim, baseline, proof packet, evals, cost/risk ledger, validation verdict, and delayed-outcome plan.	Turns proof into institutional trust.
RSI Dossier Program	Capture accepted and rejected upgrade candidates, falsification tests, rollback events, and capacity reinvestment decisions.	Prevents novelty theater and makes improvement cumulative.

9. Goals Used

Goal	Operational target	Measurement
Institutional trust	Make claims inspectable through Evidence Dockets and Proof Cards.	Independent replay, reviewer agreement, false acceptance rate.
Sovereign RSI	Make improvements cumulative without allowing unverified propagation.	Accepted artifacts beating baselines under equal budget.
Proof-settled work	Use AGIALPHA to coordinate work only when proof and validation exist.	No settlement without ProofBundle and validation.
Capacity compounding	Reinvest verified gains into stronger future work.	Useful work per cost, cost/watt, reuse rate, validator quality.
Claim discipline	Keep civilization-scale language as horizon, not present-tense claim.	Claim-boundary audits and external review.

10. Evidence Docket as Sovereign Proof Room

The Evidence Docket is the institutional proof room. It contains the claim, baseline, mission commitment, policy context, proof bundle, replay path, evaluator attestations, cost ledger, risk ledger, settlement receipt, challenge status, rollback plan, delayed-outcome monitor, and public/private boundary. Without a docket, a claim remains a story. With a docket, it becomes auditable.

Docket section	Purpose	Failure if absent
Claim matrix	States what is being claimed and what is not being claimed.	Marketing can outrun evidence.
Baseline ledger	Defines static, human, agentic, and prior-version comparators.	Improvement cannot be measured.
ProofBundle	Packages hashes, logs, traces, telemetry, signatures, replay path, and settlement evidence.	Work cannot be reproduced or disputed.
Reviewer record	Stores evaluator attestations, challenges, and decisions.	Validation becomes unaccountable.
Cost/risk ledger	Records spend, latency, risk class, safety notes, and policy status.	Utility may be overstated.
Rollout/rollback	Defines canary scope and recovery target.	Propagation becomes unsafe.


11. Settlement, Credentials, Reputation, and Access

Accepted proof does more than release rewards. It issues public-safe Proof Cards, non-transferable credentials, reputation updates, and better-mission eligibility. This turns work into trust. It also prevents unproven claims from contaminating future routing decisions.

Primitive	Function	Why it compounds
Proof Card	Public-safe summary of what improved and how it was validated.	Sponsors can understand progress without seeing private data.
Credential	Non-transferable record of approved proof work.	Builders and reviewers become legible across future missions.
Reputation	Weighted history of approved work, review quality, challenges, and reliability.	Routing improves because trust is based on proof, not marketing.
Proof-backed upgrade right	Earned right for an artifact to influence future work.	Capability compounds without unverified contagion.


12. Capacity Reinvestment and the Strategic Horizon

The civilization-scale horizon is treated as a measurement ladder, not a present-tense claim. The near-term objective is to convert verified work into reusable capability, then into productive capacity, then into compute, tools, data, validator quality, energy efficiency, and stronger future missions. This is the practical bridge from work to capacity.



13. Readiness Ladder

Architecture, evidence tooling, and repository hygiene can be strong before deployment authorization exists. Mainnet authorization remains evidence-gated: compile, tests, rehearsal, docket, external review, legal/tax/security review, governance ceremony, public claims review, and founder approval.



14. Season 001 Proof Mission Program

Season 001 should begin with a narrow proof program, not a broad-market launch. The first missions must create proof cards, reviewer history, and evidence discipline. Five missions are sufficient to establish the loop: support workflow, claims-safe launch, workshop qualification, sponsor mission design, and reviewer quality.

Season 001 Proof Mission Program

PM-001	Support Workflow	Proof Card 001
PM-002	Claims-Safe Launch	Public copy gate
PM-003	Workshop Qualification	Pipeline quality
PM-004	Sponsor Missions	Demand creation
PM-005	Reviewer Quality	Trust layer

Output discipline: every mission emits proof, reviewer verdict, claim boundary, credential decision, reputation update, and rollback note.

15. Proof Gradient Selection Law

The selection law is simple: an upgrade may influence future work only if it demonstrates positive verified value under bounded risk and rollback readiness. Formally:

$$\text{Promote}(u) = 1 \text{ iff } \text{Proof}(u) \wedge \text{EvalPass}(u) \wedge \text{RiskOK}(u) \wedge \text{ScopeOK}(u) \wedge \text{CanaryOK}(u) \wedge \text{MonitorOK}(u) \wedge \text{RollbackReady}(u).$$

Rejected candidates remain valuable. They become negative controls, failure examples, red-team material, and future benchmark cases. A proof system that hides failure cannot claim to govern improvement.

16. Falsification Tests

Failure condition	Meaning	Required response
Proof is not predictive	Accepted artifacts fail on hidden future tasks.	Tighten proof standards and reduce propagation scope.
Baselines win	Static or ordinary agent workflows beat accepted upgrades.	Reject upgrade class; publish negative result.
Gate is gameable	Candidates exploit proxy metrics or evaluator weaknesses.	Change metrics, add adversarial reviews.
Rollback fails	Promoted artifacts cannot be safely reverted.	Suspend propagation and repair rollback registry.
Evidence is unreplayable	Dockets cannot be independently inspected.	No settlement; no stronger claim.
Overhead dominates	Governance cost exceeds verified value.	Reduce coordination burden and focus on higher-value missions.

17. Claim Boundary

This paper defines a sovereign RSI architecture and a deployment program. It does not claim achieved AGI, ASI, superintelligence, autonomous sovereignty, guaranteed ROI, investment return, safety certification, legal approval, tax approval, security approval, energy abundance, or Kardashev Type II achievement. The claim becomes empirical only through real tasks, strong baselines, replayable ProofBundles, evaluator attestations, safety and cost ledgers, delayed-outcome checks, and independent review.

18. One-Page Constitution

Aim. Act. Prove. Evolve. GoalOS gives the institution direction. PlanOS gives it strategy. Capability OS gives it reusable ability. AGIALPHA coordinates proof-settled work. AEP-001 defines valid evidence. The Proof Gradient decides what may evolve. The Evolution Ledger records what happened. The Evidence Docket makes claims auditable. The Selection Gate promotes only what proved itself.

A model can answer. An agent can act. A sovereign intelligence institution must prove.

19. Formal Object Model

The architecture is implementable because every important institutional event has a typed object. The point is not more paperwork; the point is that every future upgrade can be traced back to a mission, a run, a proof packet, a validation decision, a settlement event, a rollout receipt, and a rollback target.

Object	Owner	Decision it enables
GoalOSCommit	Aim Council	What should be attempted, under whose authority, and with which constraints.
RunCommitment	Strategy Desk	Which agents, tools, context, policies, credentials, and budgets are allowed.
ProofBundle	Proof Engineers	Whether the work can be inspected, replayed, challenged, and settled.
EvalAttestation	Reviewer Council	Whether the candidate passed tests, baselines, risk gates, and policy checks.
SelectionCertificate	Selection Gate	Whether the artifact may influence future work.
RolloutReceipt	Release Office	Where the artifact may operate, at what scope, and under what monitor.
RollbackReceipt	Rollback Desk	How the system reverts or contains a failed improvement.
ChronicleEntry	Evidence Office	How the organization remembers what happened and why.

20. AGIALPHA Utility Boundary

AGIALPHA is used as a coordination asset for proof-settled work. It coordinates posting, bonding, review, settlement, credentials, reputation, and protocol operations. The token does not define truth; evidence does. The token does not replace product revenue, legal review, tax review, or security review. It coordinates utility only where public proof, accountability, and settlement are needed.

Use	Permitted role	Boundary
Mission posting	Sponsor stakes a reward and posts a bounded proof mission.	No investment framing.
Builder claim	Builder posts a bond and accepts the evidence requirements.	No entitlement to reward without accepted proof.
Proof submission	Builder submits proof references and proof-card hash.	Private evidence remains off-chain.
Reviewer validation	Reviewer stakes reputation and/or bond to validate evidence.	Review can be challenged.
Settlement	Reward and bonds settle only after validation clears.	No replay, no settlement.
Reputation	Accepted work updates non-transferable reputation.	Reputation follows proof, not token balance alone.

21. Sovereign Domain Playbook

The architecture is designed for sovereign domains where the cost of unverified action is high and the value of cumulative proof is enormous. Each domain begins with one repeatable workflow, not an unrestricted mandate. The proof room expands only when evidence and rollback readiness exist.

Domain	First proof mission	Sovereign value
AI governance	Claims-safe launch and policy review workflow.	Reduces public-claims risk and improves institutional trust.
Cybersecurity	Defensive evidence docket for safe remediation readiness.	Turns findings into governed improvements without leaking sensitive data.
Energy / compute	Useful compute per dollar and watt improvement workflow.	Links verified efficiency to reinvested compute and energy capacity.
Public services	Eligibility, routing, and citizen-response workflow.	Improves service quality while preserving auditability and rollback.
Enterprise operations	Support, procurement, and trust-center response workflow.	Creates repeatable evidence for departments and buyers.
Scientific operations	Experiment planning and validation workflow.	Improves research execution while keeping claims bounded and replayable.

22. Benchmark and Proof Program

The system should be judged by whether proof predicts future improvement, whether accepted artifacts beat baselines, whether reviewer decisions are reliable, whether rollback works, and whether coordination overhead remains below verified value. The benchmark program is adversarial by design.

Benchmark	Question	Promotion requirement
GoalOS-Bench	Can aims be specified, bounded, and audited before agents act?	Commitments are complete, authorized, and rollback-aware.
ProofGradient-Bench	Do proof packets predict hidden-task improvement?	Accepted upgrades beat strong baselines under equal budget.
SelectionGate-Bench	Can the gate reject proxy gaming and unsafe novelty?	Negative controls are rejected and failure cases are published.
EvidenceDocket-Bench	Can independent reviewers inspect and replay the evidence?	Dockets reduce audit time and false approvals.
SovereignPilot-Bench	Can one real workflow become reusable capability?	Proof Card, credential, reputation, and delayed-outcome monitor are complete.
Capacity-Bench	Does verified work improve productive capacity?	Useful work per cost/watt/time improves without hidden risk.

23. Implementation Gate Register

The implementation path is intentionally gated. Strong architecture is not deployment authorization. Each gate creates a stronger artifact, and no gate should be skipped for appearance of speed.

Gate	Evidence required	Decision
Repository readiness	Static QA, safety scan, no private data, no paid product leakage.	Engineering handoff.
Compile and tests	Install, compile, test suite, static analysis, readiness verifier.	Sepolia rehearsal permission.
Ethereum Sepolia rehearsal	Deployment manifest, mission flow, proof card, credential, reputation, settlement.	External audit closure request.
Evidence Docket	Filled docket with transaction links, proof bundle, review record, cost/risk ledger.	Claims review.
External review	Security findings, closure matrix, remediation evidence.	Governance ceremony preparation.
Governance ceremony	Admin roles, treasury, pauser, reviewer manager, founder approval.	Mainnet discussion only.
Mainnet authorization	Legal, tax, security, claims, treasury, evidence, and founder signoff.	Deployment may be considered.

24. Strategic Interpretation

The strongest form of this architecture is not a claim that the system has already achieved superintelligence. It is the claim that sovereign intelligence organizations need a way to make improvement legible, replayable, challengeable, settleable, and reversible. The civilization-scale horizon becomes credible only through a ladder of verified work: one proof room, one proof card, one benchmark, one successful rollback, one capacity gain, one validated reinvestment decision at a time.

References

Boucher, V. GoalOS: The Proof-of-Evolution Constitution AEP-001 Institutional Standard. QUEBEC.AI & MONTREAL.AI, June 2026.

Boucher, V. AGI ALPHA: A Scalable Substrate for Intelligence Organizations. QUEBEC.AI & MONTREAL.AI, June 2026.

Ethereum smart-contract documentation; Solidity documentation; OpenZeppelin contracts; Ethereum attestation and content-addressing references.

NIST AI Risk Management Framework; W3C Verifiable Credentials; OWASP agentic AI guidance.