

AGIALPHA Autopilot Safety and Claim Boundary v2

What the automation may do, what it must never do, and what claims are allowed

This pack does not claim achieved AGI, ASI, superintelligence, guaranteed ROI, yield, legal approval, tax approval, external audit completion, production safety, token price appreciation, or Ethereum Mainnet deployment. It defines a claim-bounded automation layer for proof, budgeting, review, settlement readiness, and treasury routing.

Allowed automation

- Generate mission plans, budget quotes, funding instructions, proof bundles, evidence dockets, reviewer packets, alpha-WU scoring, settlement-readiness reports, treasury routing recommendations, and public-safe evidence indexes.
- Run simulation and Sepolia rehearsal workflows.
- Open PRs with generated evidence.
- Produce mainnet read-only operator packets with no private key and no broadcast capability.

Forbidden automation

- No Mainnet broadcast from CI or GitHub Actions.
- No automatic Mainnet settlement before external audit, external evidence, and human approval.
- No slashing automation without a dispute process.
- No public claim of production safety or external audit completion without evidence.
- No token-price, ROI, yield, dividend, profit-share, or investment-return claims.
- No secret logging or private key storage in the repo.

Settlement gate

Requirement	Reason
ProofBundle complete	Settlement requires replayable evidence.
Evaluator pass	Validators gate proof quality.
Safety ledger clean	Unsafe work cannot settle as success.
Rollback path	Promotion must be reversible or stoppable.

Challenge window	Reviewer disagreement must be possible.
Claim boundary	Evidence pages must not overclaim.

Canonical slogan

No ProofBundle -> no settlement. No eval -> no propagation. No rollback -> no release. No Evidence Docket
-> no strong claim. No human gate -> no Mainnet value movement.